

# Generating trusted sphinx packets

Aurélien Chassagne  
Université Libre de Bruxelles  
Brussels, Belgium  
aurelien.chassagne@ulb.be

Jan Tobias Mühlberg  
Université Libre de Bruxelles  
Brussels, Belgium  
jan.tobias.muehlberg@ulb.be

Iness Ben Guirat  
Université Libre de Bruxelles  
Brussels, Belgium  
iness.ben.guirat@ulb.be

Jean-Michel Dricot  
Université Libre de Bruxelles  
Brussels, Belgium  
jean-michel.dricot@ulb.be

## Abstract

We propose a decentralized scheme that prevent mixnets users from sending traffic that does not match the service provider of an anonymous credential. Our scheme is of direct use in the Nym network where users construct an anonymous credential given they receive a certificate of them paying for that service provider. Our scheme prevent users from cheating by using a credential for a free service and send traffic for a paid service. Our solution works even if the majority of the third parties collude. Finally we evaluate the performance of our solution.

## CCS Concepts

• **Security and privacy** → *Privacy-preserving protocols; Distributed systems security.*

## Keywords

mixnet, sphinx, malicious users

## ACM Reference Format:

Aurélien Chassagne, Iness Ben Guirat, Jan Tobias Mühlberg, and Jean-Michel Dricot. 2025. Generating trusted sphinx packets. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (WPES '25)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/XXXXXXX.XXXXXXX>

Target: WPES - 24th Workshop on Privacy in the Electronic Society;  
deadline: 2025-07-07;  
Link: <http://jianying.space/WPES2025/>

Submission guideline:  
Submissions should be no more than 12 pages in the ACM double-column format, excluding the bibliography and well-marked appendix.  
Committee members are not required to read the appendices, and so the paper should be intelligible without them.  
Submissions should not be anonymized.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

WPES '25, Taipei, Taiwan

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 978-1-4503-XXXX-X/2018/06  
<https://doi.org/XXXXXXX.XXXXXXX>

## 1 Introduction

Mix network, or mixnet, is an overlay network of servers (called mix nodes) that prevent an adversary from correlating senders with receivers [3, 4, 6, 8, 11, 14, 16, 17]. They achieve this by adding delays to messages inside the different mixnodes in order to mitigate timing analysis attacks. Additionally, unlike Tor [10] where all packets sent by a user follow the same path for the entire session (called circuit-based), in mixnets routing is typically packet-based (meaning that each packet follows a different path). These two techniques ensure that mixnets are resilient against a strong adversary who observes the entire input and outputs of the network typically called a Global Passive Adversary (GPA). During the last few decades several systems have been proposed in the literature and few of them have been implemented. Open research problems such as parameterization, authentication and fake dummies have halted the usage on a large scale.

what are "fake dummies" do we still say "dummies" :-s

Recently, Nym Technologies a mixnet-based company is being commercialized and proposing a mixnet network for services to be integrated with their network for a fee. Their network is based on Loopix [17]. Users of these services are then allowed to use the Nym mixnet by using Nym credentials (based on the Coconut credential [18]) that they can construct after getting a certificate of paying for a specific service to use the Nym network. However, nothing prevents users from cheating who might exploit a valid Nym

Can we quantify the impact somehow? Show that this is a real problem?

credential for another service they did not pay for. This is a particular difficult problem in anonymous communication networks where the mixnodes do not know the traffic type or the final service a user is communicating with by doing layered encryption where the final IP address is only known by the last node in the path using a packer format such as Sphinx [7].

In this paper we present a scheme that creates the Sphinx header in a decentralized way based on trusted third parties while ensuring that these parties learn nothing about the destination or the path.

This sentence needs something more about the preventing cheating and reduction of overheads.

From our call 2025-06-10: "While it's computationally more complex, our schema does not allow the mixnet to be saturated by an attacker."

Would there be a risk of unauthenticated attacker-controlled traffic somehow breaking anonymity?

Our work makes the following contributions:

- ...
- ...
- We assess the impact on availability and economic viability adversaries can have in Sphinx and our solution and conclude that...

Something on prototype and availability of artifacts!

For CBT: highlight relevance of the contribution to Lightning and Nym in intro and contributions.

We highlight related work and motivation in Section 2, then we specify and justify our system model in Section 3, where we also describe the considered threat model. We then present our scheme that decentralize the creation of the Sphinx headers in Section 4 and the evaluation of our proposed solution in Section 5. Finally we conclude and discuss future work in Section 6

## 2 Motivation and Related Work

Since Chaum's seminal work on untraceable email in 1981 [3], there has been a great amount of research related to mixnets' design [1, 3–6, 11–13, 17]. However most systems that have been deployed and used come with their own system on top of the network meaning that only one type of traffic type is allowed. As beautifully stated by Dingledine et al. in [9], "Anonymity Loves Company" meaning that the more messages there are in the network the more privacy the network provide. This is also shared by Ben Guirat et al. in [2] where the authors show that blending different traffic types on top of a mixnet provide better anonymity. Let's imagine an instant messaging system where users do not tolerate latency of more than few seconds and an email app where users tolerate latency of up to 1 minute. The authors show that blending these two types of traffic do actually increase the privacy for both traffic. This only applicable in networks such as Tor or the Nym network that they offer the network for different applications to be integrated on top of the network rather than dictating which application to use the mixnet. However, certain open research problems remain open. For example how can we ensure that certain traffic are not allowed (for whatever reason and we will specify the exact reason for our work) without compromising? Tor solves the problem with having exit policy that simply drop traffic at the last node. However Tor routing is circuit-based meaning that all traffic from a specific user session follow one path and hence packets are encrypted using symmetric cryptography which is cheap, so even if traffic is dropped that is not a problem.

However, and beyond Tor, mixnets can impose substantial computational overheads due to packet based routing where every packet takes a different path and hence layered public-key cryptography is used. This implies that, if packets are dropped at the last node of the mix, which is traditionally the only place where service contracts can be validated, computational resources required by the mix to propagate a packet to its final node are wasted. Additionally, and unlike Tor where relaying traffic happens on a volunteering base, nodes in mixnets such as Nym are economically incentivized through payment per relayed bandwidth. Thus, dropping packets at an the exit node impacts the economic viability of the mixnet service.

- ...

Some concrete research questions that we can derive here and answer?

## 3 System and Threat Model

We observed that adversarial interactions on mixnets can increase the use of computational resources and network resources, thus reducing the availability and negatively affecting economic viability of mixnet services. Our work seeks to mitigate these challenges. Below we describe the system model and the adversary we consider, and derive privacy and security requirements for our solution.

### 3.1 System Model

Aurelien: I think we could try to be more concise on this subsection because it feels quite heavy (especially the second half speaking of coconut and blockchain). But for the moment, this part is a bit tricky since we haven't really focus on credentials...

Iness: add a simple graph here of a mixnet; purpose: illustrate system model, benign interactions, adversarial interactions.

We abstract the Nym mixnet to the components that are relevant for our work. We consider a mixnet, where each user have a gateway that shows a credential and if accepted their traffic is allowed. To prevent correlation, mixnet relies on fixed-size packet format such as Sphinx packet, making it difficult for external observers to link incoming and outgoing messages at any given node. The Nym credential is constructed by the user and issued by a third decentralized party after obtaining a certificate that proves payment. For example let's say Signal is integrated with Nym, and Signal users who want their traffic to be anonymous instead of sending traffic directly to the Signal server, traffic will be first routed through the mixnet such that an adversary who observes the signal server and/or the device of the user can not correlate the sender with signal server and eventually the final recipient. Signal (service provider) can add an option for user who want to pay and issue a certified attribute to those users. Users then encode this attributes into a credential and sends it to validators. If the proof is valid, validators return partial signatures. Once the user collects a threshold number of these signatures, they aggregate them to form a valid credential and re-randomize it to ensure unlinkability from previous interactions. The user can then present this credential to a verifier to prove their right to access a service to show that the credential meets all necessary payment and authentication conditions. To prevent double-spending, the verifier checks that the credential has not already been used by consulting the blockchain and then commits the credential's serial number to the blockchain upon acceptance. For example, a user can obtain an certification from the Signal service provider, construct a valid credential and then use it to route traffic to another service provider they didn't pay for or simply not allowed (an illegal website). Such misuse would be detected only at the final node of the mixnet preventing the user from accessing another application. However, prior mixnodes would have already wasted computational resources processing an invalid packet. This vulnerability enables Denial of Service (DoS) attack by exhausting mixnodes computational power with illegitimate packets.

Additionally, each encryption layer includes an integrity tag, which prevents tampering and improves the network's resistance against malicious mixnodes and active adversaries.

### 3.2 Threat Model

We consider different types of adversaries:

Aurelien: I feel like the bullet points are mixing a bit threat model (GPA, malicious user and honest-but-curious TTP) with desired properties (collusion resistant, integrity, unlinkability, ...)

- A GPA: an adversary who is able to observe all the inputs and outputs of the network. This adversary should not be able to correlate an input with an output based on the packet's appearance. This is achieved by the bitwise unlinkability of the sphinx packets.
- When constructing the headers: By using a Trusted Third Party that constructs the headers we need to make sure that even if the majority of the entities collude, no one should know the final destination of the user.
- An adversary who captures the headers is not able to change headers without intervening with the integrity check and hence mixes are able to know that the integrity check has been tampered with.
- Malicious users: Users can not cheat and create their own headers and putting the final destination different from the one they have the credential for.

### 3.3 Security & Privacy Objectives

- Cheating users: Users' traffic is only allowed to be routed if the traffic belongs to the same service provider from the credential
- even if the majority of headers issuers are colluding, they do not know which service provider the user is communicating with.
- the Sphinx headers can not be altered.
- Verifiers can verify that the headers has not been altered without revealing the service provider.
- Unlinkability between sphinx packets, the original sphinx packet that is constructed in a centralized way provide the *unlinkability* property, meaning that an adversary can not know that two packets are connected to the same user. Our scheme that decentralize the headers creation aim at providing this same property.

Some objectives may need a bit of explanation re. why they matter, for others I'm not sure how they can be evaluated.

## 4 Our Solution - Multi-Party Computation (MPC)

Our approach to ensure trust in the Sphinx header is to prevent user manipulation by decentralizing the header construction to Trusted Third Parties (TTP) through the use of Multi-Party Computation (MPC).

We consider TTPs as *honest-but-curious*. This means that they follow the protocol correctly but may attempt to infer additional information from the data they process. Our design ensures that

TTPs cannot infer any information about the shared secrets  $s_i$  or the involved mixnodes, even when TTPs collude (except one).

JT: In the Nym ecosystem, who are the TTPs, who operates them, what exactly are they trusted for?

To facilitate explanation and illustration, we describe our solution such that each packet goes through a fixed-length path consisting of three mixnodes. However, the scheme is general and can be adapted to support arbitrary path lengths.

### 4.1 Header structure

In our approach, each piece of header information is encoded as an elliptic curve point. To achieve this, the encoded string is divided into fixed-size chunks, as illustrated in Figure 2, where each chunk is a single EC point. For example, in the case of a path consisting of three mixnodes, the resulting header contains seven EC points: one destination, three mixnodes and three integrity tags.

Therefore a method to encode and decode information to and from elliptic curve points is required. Specifically, we require a mapping from integers to curve points, noted  $\text{Map}()$ , and the inverse mapping from curve points to integers, noted  $\text{Map}^{-1}()$ , such that  $\text{Map}^{-1}(\text{Map}(x)) = x$  where  $x$  is an integer.

Traditional methods achieve this by directly mapping integers to the x-coordinates of points on the curve. However, this approach could leaks information by introducing bias since nearby integers result in nearby points.

To address this, we adopt Elligator, which offers stronger privacy guarantees. Unlike the traditional approach, Elligator produces a uniformly distributed output which is computationally indistinguishable from truly random curve points. This uniformity is in line with our objectives of preserving anonymity and avoiding unlinkability.


### 4.2 Protocol description

The overall decentralized scheme is illustrated in Figure 1. (1) The client first computes a sequence of shared secrets and then splits these secrets, along with the IP addresses of the mixnodes in the path and the final destination, into  $m$  shares. (2) Each set of shares is sent to a different TTP, along with the necessary cryptographic element ( $\alpha$ ). (3) Each TTP independently computes a *partial* Sphinx header using the received shares. (4) The client then aggregates these partial headers to reconstruct the final header, (5) ready for transmission through the mixnet.

In summary, the protocol is divided into four main steps:

- **Setup:** fixes random points as generators (once but should be refreshed);
- **Client:** splits and shares the necessary information to TTPs;
- **TTP:** encrypts the routing information into a partial header;
- **Mixnode:** decrypts and forwards the header.

**4.2.1 Setup.** In our protocol, routing information is divided into seven chunks, each one encoded into an elliptic curve point (see Figure 2). These points will later be encrypted by adding a masking point of the form  $P = sG$ , where  $s$  is a shared secret and  $G$  is the base generator of the curve. However, using the same masking point  $P$  for all the seven chunks compromises the *unlinkability* property (see security note in the TTP step 4.2.3).



**Figure 1: Overview of the decentralized scheme**

To mitigate this, we use a set of seven *independent* generators  $G_j$ , one for each chunk. By *independent*, we mean that the scalar relationship between any two generators is unknown. Therefore no one knows the scalars  $z_j \in \mathbb{Z}_N$  such that  $G_j = z_j G$ .

In principle, this setup phase can be executed only once. However, to preserve unlinkability, the set of generators should be refreshed periodically. Every entity (clients, mixnodes, and TTPs) must run the setup phase to ensure they use the same set of generators.

To achieve synchronized and deterministic generators across the system, we derive them from a common seed. This seed is computed as the hash of the current timestamp, truncated to the chosen refresh interval. To produce *independent* generators, we simply increment the seed since Elligator already provides a uniform (i.e. pseudo-random) integers to points mapping.

A subtlety with Elligator is that it can map integers to points outside of the base generator's subgroup. To ensure that points remain in the generator subgroup, an extra step is required to *clear the cofactor* by multiplying the resulting point by the cofactor (Curve25519's cofactor is 8).

The final formula for generating the  $j^{\text{th}}$  chunk's generator is:

$$G_j = 8 \cdot \text{Map}(h(\text{timestamp}) + j), \quad \text{for } j = 1, \dots, 7 \quad (1)$$

**4.2.2 Client.** To send a message to a specific destination, the client first randomly chooses a path through the mixnet and generates a nonce. Using this nonce and the public keys of each mixnode along the path, the client derives a chain of shared secret integers ( $s_1, s_2, s_3$ ) as follows:

$$(s_1, s_2, s_3) \begin{cases} \alpha_i = x_i G \\ S_i = x_i K_i \\ s_i = \text{Map}^{-1}(S_i) \\ r_i = h(\alpha_{i,y} \parallel S_{i,y}) \\ x_{i+1} = x_i r_i \pmod{n} \end{cases} \quad (2)$$

where  $x_1$  is the client's nonce,  $K_i$  is the public key of the  $i^{\text{th}}$  mixnode in the path ( $K_i = k_i G$ ),  $n$  is the subgroup order (i.e.  $(n+1)G = G$ ),  $\parallel$  is the concatenation operator and subscript  $y$  refers to the  $y$ -coordinate of the point. The scalar  $r_i$  acts as a pseudo-randomizer to update the secret  $x_i$  for the next hop, ensuring that each cryptographic element  $\alpha_i$  remains unlinkable across mixnodes.

The IP addresses of the mixnodes and the destination are padded with random bits then mapped to elliptic curve points using Elligator, allowing later recovery of these addresses. Both the resulting IP points and the shared secrets are split into shares and distributed to  $m$  different TTPs. In addition, the client send the hash of each shared secret to these TTPs.

Each TTP then independently computes a partial header from the received shares and returns it to the client. The client aggregates these partial headers by performing chunk-wise point additions to reconstruct the complete encrypted header. Finally, the client sends the header along with  $\alpha_1$  to the first mixnode from the path.

**4.2.3 TTP.** Each TTP receives from the client:

- a share of the destination address (EC point noted  $\Delta$ );
- a share of each mixnode address in the path (EC points noted  $N_i$ );
- a share of each shared secrets (integers noted  $s_i$ );
- the hash of each shared secrets (integers noted  $\sigma_i$ ).

The TTP constructs a partial encrypted header by layering routing information in reverse path order, starting from the last mixnode, as illustrated by Figure 2.

At each layer, two new chunks are introduced: the mixnode's address ( $N_i$ ) and the integrity tag ( $\gamma_i$ ). To preserve a fixed-length header while allowing mixnodes to reverse the operation (as in the original design), four additional *filler chunks* are appended. These filler chunks are carefully computed such that at each layer, after applying the corresponding masking points ( $s_i G_i$ ), the last two chunks cancel out and become identity point (i.e. point at infinity). This mechanism allows these two chunks to be safely truncated at each layer, ensuring a consistent header size while preserving reversibility for mixnodes.

To initialize the header, the TTP encrypts the destination point using the shared secret of the last mixnode and append the required filler chunks. Let  $\Delta$  be the partial destination point,  $s_i$  the partial shared secret corresponding to the  $i^{\text{th}}$  mixnode in the path, and  $G_j$  the  $j^{\text{th}}$  chunk's generator. The chunks for the last layer (assuming a 3-hop path) are computed as follows:

$$\beta_3 \begin{cases} \beta_{31} = \Delta + s_3 G_1 \\ \beta_{32} = -(s_2 G_4 + s_1 G_6) \\ \beta_{33} = -(s_2 G_5 + s_1 G_7) \\ \beta_{34} = -s_2 G_6 \\ \beta_{35} = -s_2 G_7 \end{cases} \quad (3)$$

Subsequent layers ( $i = 2$  and  $i = 1$ ) are computed with Eq. 4 where  $N_{i+1}$  represents the elliptic curve point corresponding to the  $(i+1)^{\text{th}}$  mixnode in the path and  $\gamma_{i+1}$  is the integrity tag from previous layer.


$$\beta_i \begin{cases} \beta_{i1} = N_{i+1} + s_i G_1 \\ \beta_{i2} = \gamma_{i+1} + s_i G_2 \\ \beta_{i3} = \beta_{i+1} + s_i G_3 \\ \beta_{i4} = \beta_{i+12} + s_i G_4 \\ \beta_{i5} = \beta_{i+13} + s_i G_5 \end{cases} \quad (4)$$

$\beta_{i+11}$  is a light notation for  $\beta_{(i+1,1)}$ , should I put the full notation for clarity ?

For each layer, the integrity tag is computed with Eq. 5 where  $h^{(j)}(\sigma_i)$  means hashing  $j$  times  $\sigma_i$  which is the hash of the  $i^{th}$  mixnode's complete shared secret. The integrity tag consists in a pseudo-random weighted sum of the chunks, offset by a secret point derived from the shared secret integer. The pseudo-random weights (i.e.  $h^{(j)}(\sigma_i)$ ) ensure binding, preventing adversaries from modifying individual chunks. Without these weights, an attacker could add a random point  $P$  to one chunk and subtract  $P$  from another, thereby preserving the overall sum and bypassing the integrity check. The secret offset (i.e.  $s_i G$ ) prevents curious TTPs from correlating integrity tags. Since TTPs are aware of these pseudo-random weights, omitting this offset would allow them to potentially link some incoming and outgoing mixnode's packets by simply verifying integrity tags.

$$y_i = s_i G + \sum_{j=1}^5 (h^{(j)}(\sigma_i) \beta_{ij}) \quad (5)$$

When the first layer is finally computed, it is appended with the integrity tag  $\gamma_1$  and the first mixnode's point  $N_1$ . Finally, the TTP returns the constructed partial header to the client.



**Figure 2: Chunk-wise header encryption (TTP side), where (x) refers to Eq. x.**

**Security Note.** To preserve unlinkability, it is essential that the generators  $(G_1, \dots, G_7)$  are *independent*, meaning that their scalar relationships are unknown and cannot be derived. If the same generator  $G$  were used across chunks (or their scalar relationships known), an adversary could compute the chunk-wise difference between consecutive layers:  $\beta_{ij} - \beta_{i,j+1} = s_i G$ . Since the shared secret  $s_i$  remains consistent within a layer, the resulting differences would reveal a predictable pattern (uniform or preserved scalar relationships). This consistency could allow an adversary to correlate incoming and outgoing packets at a mixnode, thereby breaking

the unlinkability property. Using independent generators for each chunk prevents such correlations and is therefore critical to the protocol’s security.

**4.2.4 Mixnode.** When a packet arrives at mixnode  $i$ , it begins by extracting the relevant header fields, namely the cryptographic element  $\alpha_i$  and the integrity tag  $\gamma_i$ . The mixnode then recomputes the shared secret point  $S_i$  using its private key  $k_i$  and the cryptographic element  $\alpha_i$ :

$$\begin{aligned} S_i &= k_i \alpha_i \\ s_i &= \text{Map}^{-1}(S_i) \end{aligned} \quad (6)$$

This shared secret point is subsequently mapped to the shared integer  $s_i$  using Elligator inverse mapping. This shared integer is then used to verify the integrity tag (Eq. 7).

$$y_i = s_i G + \sum_{j=1}^5 (h^{(j)}(s_i) \beta_{ij}) \quad (7)$$

If the integrity check passes, the mixnode pads the header by appending two *identity* chunks (i.e. point at infinity) and then updates it as:

$$\beta'_{ij} = \beta_{ij} - s_i G_j \quad \forall j = 1, \dots, 7 \quad (8)$$


The first chunk ( $\beta'_{i1}$ ), encodes the next mixnode’s IP address ( $N_{i+1}$ ). This point is mapped back to an integer using Elligator and the random padding is removed by keeping the last 128 bits (i.e. size of an IP address).

The second chunk ( $\beta'_{i+1}$ ) is the next integrity tag ( $y_{i+1}$ ) for the remaining five chunks that form the new encrypted routing information ( $\beta_{i+1}$ ).

To maintain unlinkability, the cryptographic element  $\alpha_i$  must be updated for the next hop. As in the client's step, it is computed using the  $y$ -coordinates of both  $\alpha_i$  and  $S_i$ :

$$\alpha_{i+1} = \text{hash}(\alpha_{i_y} \parallel s_{i_y}) \alpha_i \quad (9)$$

Finally, the mixnode forwards the updated header  $(\alpha_{i+1}, \gamma_{i+1}, \beta_{i+1})$  to the next node.



**Figure 3: Processing sphinx header at mixnode  $i$ , where  $\textcircled{x}$  refers to Eq. x.**

## 5 Evaluation

Our prototype implementation is available on GitHub<sup>1</sup>. The implementation is written in Python 3.13 using the ECPy library

<sup>1</sup><https://github.com/AurelienCha/Decentralized-Sphinx>

**Table 1: Number of expensive cryptographic operations per party, where  $p$  is the path length (e.g.  $p = 3$ ), and  $m$  is the number of TTPs.**

|                             | Client   | TTP      | Mixnode |
|-----------------------------|----------|----------|---------|
| EC Multiplication           | $O(p m)$ | $O(p^2)$ | $O(p)$  |
| EC Addition                 | $O(p m)$ | $O(p^2)$ | $O(p)$  |
| Point $\rightarrow$ Integer | $O(p)$   | 0        | 2       |
| Integer $\rightarrow$ Point | $O(p)$   | 0        | 0       |
| Hash                        | $O(p)$   | $O(p^2)$ | $O(p)$  |

for elliptic curve operations. ECPy is well-suited for prototyping due to its simplicity, but it is significantly slower than many EC python libraries (approximately one order of magnitude slower). Nevertheless, it remains one of the few Python libraries supporting Montgomery Curve25519, which is essential for our implementation.

We evaluate our system on two main axes: computational complexity and unlinkability (i.e. resistance to correlation between incoming and outgoing packets).

## 5.1 Computational Complexity

Instead of measuring raw execution time, which can vary significantly based on hardware and software environments, we focus on the number of expensive cryptographic operations. Table 1 count these operations per party for  $m$  TTPs and path of length  $p$ . To compute the total cost for sending one message, the TTP cost should be multiplied by  $m$  and the mixnode cost by  $p$ .

Since long paths are typically unnecessary, we can approximate  $O(p)$  as constant (i.e.  $O(1)$ ). A path of length  $p = 3$  is generally sufficient to ensure strong anonymity.


Aurelien: Iness do you have some source to support this ? If yes, could you cite here, thanks.

Consequently, the overall computational burden is more sensitive to the number of TTPs  $m$  rather than the path length. Since our design remains secure as long as a single TTP is honest (i.e. does not collude), only a small number of TTPs may be required. A deeper analysis of how many TTPs are needed to ensure a desired level of trust remains an open question for future work.

## 5.2 Unlinkability assessment

Quantifying unlinkability is inherently challenging. However, under the assumption that uniformly random packet headers prevent adversaries from correlating incoming and outgoing packets through cryptanalysis, unlinkability can be approximated by assessing the statistical randomness of the headers.

For this purpose, we rely on the NIST SP 800-22 statistical test suite [15], which consists of 15 tests designed to evaluate the quality of random number generators. Each test produces a p-value representing how likely the data could be produced by a uniform source. If the data are truly random, the distribution of p-values across many samples should itself follow a uniform distribution.



**Figure 4: Proportion of successful sequences from our implementation and the original Sphinx implementation over the fifteen NIST SP 800-22 statistical tests (significance level = 0.01).**

We compare the statistical results of our implementation with those of the original Sphinx implementation by Danezis<sup>2</sup>. All graphs presented in this section correspond to the results of the NIST statistical test suite applied to 20 distinct datasets. Each dataset is generated under a different setup configuration to assess the variability and robustness of our implementation.

To generate a dataset, we simulate a network of 20 mixnodes and 3 trusted third parties (TTPs). For each iteration, a random destination IP address and a path consisting of 3 mixnodes are selected. The client then generates shares and distributes them to the three TTPs, which independently compute partial headers. These headers are then aggregated and forwarded through the simulated mixnet.

According to the NIST documentation, the statistical test suite requires input sequences of 1 000 000 bits. To meet this requirement, generated headers are concatenated until the total length reaches 1 000 000 bits. Each dataset contains 100 of those concatenated sequences.

NIST outlines two approaches for interpreting the empirical results of its statistical test suite. The first assesses the proportion of sequences that pass each statistical test, as shown in Fig.4. The second evaluates whether the distribution of p-values is statistically uniform, as would be expected from random sequences. These results are presented in Fig.6.

Figure 4 shows the proportion of sequences that successfully pass each statistical test for both our implementation and the original Sphinx implementation. Following the NIST guidelines, the significance level is set to  $\alpha = 0.01$ , and the confidence interval is defined as:

$$(1 - \alpha) \pm 3\sqrt{\frac{\alpha(1 - \alpha)}{m}},$$

where  $m$  denotes the number of samples.

<sup>2</sup><https://github.com/UCL-InfoSec/sphinx>




**Figure 5: P-value distribution comparison between our implementation and the original Sphinx across the fifteen NIST SP 800-22 statistical tests.**

Our implementation fails only three tests: *Nonperiodic Template Matching*, *Random Excursions*, and *Random Excursions Variant*, all of which are also failed by the original Sphinx implementation. In contrast, our implementation successfully passes three tests that the original implementation fails: *Frequency*, *Cumulative Sums*, and *Overlapping Template Matching*.

Figure 5 presents the distribution of p-values obtained for each test. This figure is intended solely for visual inspection. A Chi-square test is performed on each histogram to assess their uniformity. The results are shown in Figure 6, with a significance threshold of 0.0001 as recommended by the NIST guidelines.

Two Chi-square results stand out: the *Frequency* and *Cumulative Sums* tests. Both were successful for our implementation, despite failing for the original Sphinx implementation.



**Figure 6: Chi-square test assessing the uniformity of p-value distributions for each dataset across the fifteen NIST SP 800-22 statistical tests.**

### 5.3 Discussion

In terms of computational complexity, our implementation is slower than the original Sphinx implementation. However, in the context of a mixnet where inherent latency is already expected, this additional delay is likely acceptable.

**Table 2: Statistical test summary for both implementations (● = Pass, ○ = Fail)**

| Statistical Test            | Decentralized        |                  | Original             |                  |
|-----------------------------|----------------------|------------------|----------------------|------------------|
|                             | Excessive Rejections | Lacks Uniformity | Excessive Rejections | Lacks Uniformity |
| Frequency                   | ●                    | ●                | ○                    | ○                |
| Block Frequency             | ●                    | ●                | ●                    | ●                |
| Cumulative Sums             | ●                    | ●                | ○                    | ○                |
| Runs                        | ●                    | ●                | ●                    | ●                |
| Longest Run                 | ●                    | ●                | ●                    | ●                |
| Rank                        | ●                    | ●                | ●                    | ●                |
| Discrete Fourier Transform  | ●                    | ●                | ●                    | ●                |
| Nonperiodic Template Match. | ○                    | ●                | ○                    | ●                |
| Overlapping Template Match. | ●                    | ●                | ○                    | ●                |
| Universal Statistical       | ●                    | ●                | ●                    | ●                |
| Approximate Entropy         | ●                    | ●                | ●                    | ●                |
| Random Excursions           | ○                    | ●                | ○                    | ●                |
| Random Excursions Variant   | ○                    | ●                | ○                    | ●                |
| Serial                      | ●                    | ●                | ●                    | ●                |
| Linear Complexity           | ●                    | ●                | ●                    | ●                |

Regarding unlinkability (i.e. output randomness), our implementation performs slightly better than the original implementation, as it passes more statistical tests, as shown in Table 2. Since the original Sphinx implementation is widely considered as secure, and assuming no specific vulnerabilities exist in our design, we can reasonably infer that our approach maintains a comparable level of security.

Come back to 3.3 objectives and explain why they hold. If we have research questions from 2, answer them. Summarize overheads impacts and discuss feasibility in the context of Nym.

Aurelien: It is quite difficult to refer to these objectives because the paper in this state does "half" of the work. We only decentralized the construction of the Sphinx header. After that, we still need to provide the actual "credentials". It's just way easier to do it now since we can "verify/link" the credential with "partial destination" (with the original implementation we would have to somehow verify/link this credential to a destination that have been encrypted under multiple layer.)

We have proposed a decentralized scheme that delegates the construction of the Sphinx header to a set of Trusted Third Parties (TTPs), ensuring that no single TTP can infer meaningful information from the computation, even in the case of collusion (if at least one party remains honest). Since these TTPs directly access partial destination IP information, it becomes significantly easier to issue anonymous credentials for specific usages, while preventing users from cheating on their intended destinations. The detailed construction of such credentials, lies beyond the scope of this paper.

Another future work is to evaluate the computational overhead introduced by the TTPs in comparison to the computational waste caused by malicious users in the original Sphinx implementation, in order to assess the overall trade-off.

## 6 Conclusion

## Acknowledgments

This research is funded by Belgium Walloon region CyberExcellence program (grant #2110186)

## References

- [1] Nikolaos Alexopoulos, Aggelos Kiayias, Riivo Talviste, and Thomas Zacharias. 2017. MCMix: Anonymous Messaging via Secure Multiparty Computation. In *26th {USENIX} Security Symposium ({USENIX} Security 17)*. 1217–1234. <https://www.usenix.org/conference/usenixsecurity17/technical-sessions/presentation/alexopoulos>
- [2] Iness BEN GUIRAT, Debajyoti Das, and Claudia Diaz. 2023. Blending different latency traffic with beta mixing. *Proceedings on Privacy Enhancing Technologies* (2023).
- [3] David Chaum. 1981. Untraceable Electronic Mail, Return Addresses, and Digital Pseudonyms. *Commun. ACM* 24, 2 (1981), 84–88. doi:10.1145/358549.358563
- [4] David Chaum, Farid Javani, Aniket Kate, Anna Krasnova, Joeri Ruiter, Alan T Sherman, and Debajyoti Das. 2016. *cMix: Anonymization by high-performance scalable mixing*. Technical Report. Technical report. <http://www.cs.bham.ac.uk/~deruitej/papers/cmixon.pdf>
- [5] Lance Cottrell. 1995. Mixmaster and remailer attacks. <https://www.obscura.com/~loki/remailer-essay.html>
- [6] George Danezis, Roger Dingledine, and Nick Mathewson. 2003. Mixminion: Design of a Type III Anonymous Remailer Protocol. In *Proceedings of the 2003 IEEE Symposium on Security and Privacy*. IEEE, 2–15.
- [7] George Danezis and Ian Goldberg. 2009. Sphinx: A Compact and Provably Secure Mix Format. In *2009 30th IEEE Symposium on Security and Privacy*. 269–282. <http://research.microsoft.com/en-us/um/people/gdane/papers/sphinx-eprint.pdf>
- [8] Claudia Diaz, Harry Halpin, and Aggelos Kiayias. 2021. The Nym Network. <https://nymtech.net/nym-whitepaper.pdf>, 38 pages.
- [9] Roger Dingledine and Nick Mathewson. 2006. Anonymity Loves Company: Usability and the Network Effect.. In *WEIS*.
- [10] David M. Goldschlag, Michael G. Reed, and Paul F. Syverson. 1996. Hiding Routing Information. In *Proceedings of the First International Workshop on Information Hiding*. Springer-Verlag, Berlin, Heidelberg, 137–150.
- [11] Jelle Van Den Hooff, David Lazar, Matei Zaharia, and Nickolai Zeldovich. 2015. Vuvuzela: Scalable private messaging resistant to traffic analysis. In *Proceedings of the 25th Symposium on Operating Systems Principles*. 137–152. <https://people.csail.mit.edu/nickolai/papers/vandenhooff-vuvuzela.pdf>
- [12] Albert Kwon, David Lu, and Srinivas Devadas. 2020. {XRD}: Scalable messaging system with cryptographic privacy. In *17th {USENIX} Symposium on Networked Systems Design and Implementation ({NSDI} 20)*. 759–776.
- [13] David Lazar, Yossi Gilad, and Nickolai Zeldovich. 2018. Karaoke: Distributed private messaging immune to passive traffic analysis. In *13th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 18)*. 711–725.
- [14] Ulf Möller, Lance Cottrell, Peter Palfrader, and Len Sassaman. 2003. Mixmaster Protocol — Version 2. IETF Internet Draft.
- [15] NIST. 2010. *A Statistical Test Suite for Random and Pseudorandom Number Generators for Cryptographic Applications*. Special Publication SP 800-22. National Institute of Standards and Technology. doi:10.6028/NIST.SP.800-22r1a
- [16] Sameer Parekh. 1996. Prospects for remailers. *First Monday* 1 (August 1996). Issue 2.
- [17] Ania M Piotrowska, Jamie Hayes, Tariq Elahi, Sebastian Meiser, and George Danezis. 2017. The loopix anonymity system. In *26th {USENIX} Security Symposium ({USENIX} Security 17)*. 1199–1216.
- [18] Alberto Sonnino, Mustafa Al-Bassam, Shehar Bano, Sarah Meiklejohn, and George Danezis. 2019. Coconut: Threshold issuance selective disclosure credentials with applications to distributed ledgers. In *26th Annual Network and Distributed System Security Symposium, NDSS 2019*. The Internet Society.

Received 27 June 2025; revised XX July 2025; accepted XX July 2025